

dwl_final

May 5, 2021

0.0.1 I created a scraping script using libraries requests and BeautifulSoup. Also I created empty lists and wrote down all the information there. When we go to the csv file, we will see 10 columns and 600 lines in it. It contains following information about each car: Car, Price, Year, Region, Body, Engine_volume, Mileage, Transmission, Steering_wheel, Color.

```
[11]: import requests
import re
from bs4 import BeautifulSoup
import pandas as pd

car = []
price = []
year_of_manufacture = []
region = []
body = []
engine_volume = []
mileage = []
transmission = []
steering_wheel = []
color = []

headers = {'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/
↳537.36 (KHTML, like Gecko) Chrome/90.0.4430.93 Safari/537.36'}
for i in range(30):
    results = requests.get(f'https://kolesa.kz/cars/avtomobili-s-probegom/?
↳auto-car-order=1&auto-car-volume[from]=1&auto-car-volume[to]=9&page={i}')
    soup = BeautifulSoup(results.text, 'html.parser')
    for j in soup.find_all('span', class_='a-el-info-title'):
        car.append(j.text.strip())
        results2 = requests.get('https://kolesa.kz' + j.
↳find('a', class_='list-link ddl_product_link')['href'])
        soup2 = BeautifulSoup(results2.text, 'html.parser')

        p = re.sub("[^0-9]", "", soup2.find('div', class_='offer__price').text.
↳replace(u'\xa0', ' '))
        price.append(int(p))
```

```

        year_of_manufacture.append(int(soup2.find('span', class_='year').text.
→strip()))
        region.append(soup2.find_all('dd',class_='value')[0].text.strip())
        body.append(soup2.find_all('dd',class_='value')[1].text.strip())

        engine_volume.append(re.compile(r'[^\.d.]+').sub('.', (soup2.
→find_all('dd', class_='value')[2].text)))

        m = soup2.find_all('dd',class_='value')[3].text.strip()
        if m[-1] == ' ':
            m2 = re.sub("[^0-9]", "", m)
            mileage.append(int(m2))
            transmission.append(soup2.find_all('dd',class_='value')[4].text.
→strip())
            steering_wheel.append(soup2.find_all('dd',class_='value')[5].text.
→strip())
            color.append(soup2.find_all('dd',class_='value')[6].text.strip())
        else:
            mileage.append(0)
            transmission.append(soup2.find_all('dd',class_='value')[3].text.
→strip())
            steering_wheel.append(soup2.find_all('dd',class_='value')[4].text.
→strip())
            color.append(soup2.find_all('dd',class_='value')[5].text.strip())

myDict = {}

myDict['Car'] = car
myDict['Price'] = price
myDict['Year'] = year_of_manufacture
myDict['Region'] = region
myDict['Body'] = body
myDict['Engine_volume'] = engine_volume
myDict['Mileage'] = mileage
myDict['Transmission'] = transmission
myDict['Steering_wheel'] = steering_wheel
myDict['Color'] = color

df = pd.DataFrame.from_dict(myDict)
df.to_excel('kolesa2.xlsx', encoding='utf-8-sig')
df.to_csv('kolesa.csv', encoding='utf-8-sig')

df

```

```

[11]:
      Car      Price  Year      Region \
0      Chevrolet Tracker  5850000  2015
1          Opel Frontera  1650000  1994
2  Renault Sandero Stepway  3990000  2014
3      Toyota Land Cruiser 16800000  2010
4      Mitsubishi Pajero  8000000  2007
..      ...      ...      ...      ...
595      Daewoo Nexia  950000  2007
596      6400000  2015
597      BMW X5  7000000  2006  -  (  )
598  Volkswagen Transporter  1600000  1991
599  Mercedes-Benz E 230  1800000  1992  -  (  )

      Body Engine_volume  Mileage Transmission Steering_wheel \
0      1.8      46000
1      2.3      0
2      1.6  110000
3      4  220000
4      3  137000
..      ...      ...      ...      ...
595      1.5  190000
596      2.9      0
597      4.8  220000
598      2  222222
599      2.3      222

      Color
0
1
2
3
4
..      ...
595
596
597
598
599

[600 rows x 10 columns]

```

0.0.2 Return DataFrame with duplicate rows removed and remove null properties.

```

[52]: df = pd.read_csv('kolesa.csv')
df = df.drop('Unnamed: 0', axis='columns')
df= df.dropna()
df= df.drop_duplicates()

```

```
df.head()
```

```
[52]:
```

| | Car | Price | Year | Region | Body \ |
|---|-------------------------|----------|------|--------|--------|
| 0 | Chevrolet Tracker | 5850000 | 2015 | | |
| 1 | Opel Frontera | 1650000 | 1994 | | |
| 2 | Renault Sandero Stepway | 3990000 | 2014 | | |
| 3 | Toyota Land Cruiser | 16800000 | 2010 | | |
| 4 | Mitsubishi Pajero | 8000000 | 2007 | | |

| | Engine_volume | Mileage | Transmission | Steering_wheel | Color |
|---|---------------|---------|--------------|----------------|-------|
| 0 | 1.8 | 46000 | | | |
| 1 | 2.3 | 0 | | | |
| 2 | 1.6 | 110000 | | | |
| 3 | 4.0 | 220000 | | | |
| 4 | 3.0 | 137000 | | | |

```
[53]: df.sort_values(by=['Price'])
```

```
[53]:
```

| | Car | Price | Year | Region | \ |
|-----|---------------------------|----------|------|--------|-----|
| 192 | (Lada) 2107 | 350000 | 2006 | | |
| 79 | (Lada) 2114 () | 380000 | 2003 | | |
| 238 | (Lada) 2107 | 400000 | 2008 | | |
| 564 | 968 | 500000 | 1989 | | |
| 118 | Opel Omega | 550000 | 1987 | | |
| .. | ... | ... | ... | | |
| 412 | Toyota Land Cruiser | 31600000 | 2017 | - | () |
| 257 | Lexus RX 350 | 33500000 | 2019 | | |
| 228 | Toyota Land Cruiser Prado | 36200000 | 2021 | | |
| 55 | BMW X5 | 44300000 | 2019 | | |
| 578 | Land Rover Range Rover | 55000000 | 2018 | - | () |

| | Body | Engine_volume | Mileage | Transmission | Steering_wheel | \ |
|-----|------|---------------|---------|--------------|----------------|---|
| 192 | | 1.5 | 0 | | | |
| 79 | | 1.6 | 280000 | | | |
| 238 | | 1.7 | 0 | | | |
| 564 | | 1.2 | 80000 | | | |
| 118 | | 2.0 | 0 | | | |
| .. | ... | ... | ... | ... | ... | |
| 412 | | 4.6 | 53000 | | | |
| 257 | | 3.5 | 9000 | | | |
| 228 | | 4.0 | 0 | | | |
| 55 | | 3.0 | 40 | | | |
| 578 | | 3.0 | 25500 | | | |

| | Color |
|-----|-------|
| 192 | |
| 79 | |

238
564
118
..
412
257
228
55
578

[517 rows x 10 columns]

```
[54]: df.groupby(['Year']).first()
```

```
[54]:
```

| | Car | Price | Region | Body \ |
|------|---------------------------|----------|--------|--------|
| Year | | | | |
| 1987 | Opel Omega | 550000 | | |
| 1989 | Mazda 626 | 1450000 | | |
| 1990 | Mazda 626 | 850000 | | |
| 1991 | Opel Vectra | 950000 | | |
| 1992 | Toyota Land Cruiser Prado | 6800000 | | |
| 1993 | Volkswagen Passat | 1550000 | | |
| 1994 | Opel Frontera | 1650000 | | |
| 1995 | Mitsubishi RVR | 1600000 | | |
| 1996 | Toyota Caldina | 2250000 | - | |
| 1997 | Toyota Mark II | 2550000 | | |
| 1998 | Mercedes-Benz E 280 | 2870000 | | |
| 1999 | Mazda MPV | 3000000 | | |
| 2000 | Daewoo Matiz | 900000 | | |
| 2001 | Opel Zafira | 2300000 | | |
| 2002 | Lexus ES 300 | 5950000 | | |
| 2003 | Toyota Camry | 3800000 | | |
| 2004 | Land Rover Range Rover | 3000000 | | |
| 2005 | Land Rover Range Rover | 5800000 | | |
| 2006 | Toyota Land Cruiser | 7100000 | | |
| 2007 | Mitsubishi Pajero | 8000000 | | |
| 2008 | Mitsubishi Pajero | 7500000 | | |
| 2009 | Daewoo Nexia | 1350000 | | |
| 2010 | Toyota Land Cruiser | 16800000 | | |
| 2011 | Hyundai Sonata | 3700000 | | |
| 2012 | (Lada) 2190 () | 2190000 | - | () |
| 2013 | Nissan Juke | 5950000 | | |
| 2014 | Renault Sandero Stepway | 3990000 | | |
| 2015 | Chevrolet Tracker | 5850000 | | |
| 2016 | Chevrolet Spark | 3500000 | | |
| 2017 | Toyota Tundra | 23000000 | | |
| 2018 | Hyundai Tucson | 11350000 | | |

| | | |
|------|---------------------------|----------|
| 2019 | Toyota Camry | 13900000 |
| 2020 | Toyota Camry | 15800000 |
| 2021 | Toyota Land Cruiser Prado | 36200000 |

| | Engine_volume | Mileage | Transmission | Steering_wheel | Color |
|------|---------------|---------|--------------|----------------|-------|
| Year | | | | | |
| 1987 | 2.0 | 0 | | | |
| 1989 | 2.0 | 310915 | | | |
| 1990 | 2.0 | 0 | | | |
| 1991 | 2.0 | 145000 | | | |
| 1992 | 4.0 | 100000 | | | |
| 1993 | 1.8 | 250000 | | | |
| 1994 | 2.3 | 0 | | | |
| 1995 | 2.0 | 444444 | | | |
| 1996 | 1.8 | 0 | | | |
| 1997 | 2.0 | 0 | | | |
| 1998 | 2.8 | 0 | | | |
| 1999 | 2.5 | 0 | | | |
| 2000 | 8.0 | 0 | | | |
| 2001 | 2.2 | 160000 | | | |
| 2002 | 3.0 | 154000 | | | |
| 2003 | 2.4 | 0 | | | |
| 2004 | 4.4 | 0 | | | |
| 2005 | 4.4 | 165000 | | | |
| 2006 | 4.7 | 182000 | | | |
| 2007 | 3.0 | 137000 | | | |
| 2008 | 3.0 | 219000 | | | |
| 2009 | 1.6 | 126000 | | | |
| 2010 | 4.0 | 220000 | | | |
| 2011 | 2.4 | 197102 | | | |
| 2012 | 1.6 | 146000 | | | |
| 2013 | 1.6 | 39000 | | | |
| 2014 | 1.6 | 110000 | | | |
| 2015 | 1.8 | 46000 | | | |
| 2016 | 1.0 | 135000 | | | |
| 2017 | 4.6 | 0 | | | |
| 2018 | 2.0 | 53000 | | | |
| 2019 | 2.5 | 18000 | | | |
| 2020 | 2.5 | 1955 | | | |
| 2021 | 4.0 | 0 | | | |

```
[55]: df.groupby(['Year']).get_group(2008)
```

```
[55]:
```

| | Car | Price | Year | Region \ |
|----|-------------------|---------|------|----------|
| 7 | Mitsubishi Pajero | 7500000 | 2008 | |
| 14 | Toyota Highlander | 8800000 | 2008 | |
| 41 | Chrysler 300C | 5500000 | 2008 | |

| | | | | |
|-----|---------------------------|----------|------|-------|
| 78 | | 3199000 | 2008 | |
| 151 | (Lada) 2114 () | 1000000 | 2008 | |
| 182 | Toyota Camry | 4000000 | 2008 | |
| 185 | Toyota Camry | 4000000 | 2008 | |
| 206 | (Lada) 2113 () | 550000 | 2008 | |
| 217 | Hyundai Accent | 2100000 | 2008 | |
| 222 | BMW X6 | 7500000 | 2008 | |
| 225 | Daewoo Nexia | 550000 | 2008 | - () |
| 229 | (Lada) 2114 () | 650000 | 2008 | |
| 238 | (Lada) 2107 | 400000 | 2008 | |
| 250 | Hyundai Tuscani | 3100000 | 2008 | |
| 253 | (Lada) 2114 () | 900000 | 2008 | |
| 356 | | 4400000 | 2008 | - () |
| 361 | Volkswagen Golf Plus | 3300000 | 2008 | |
| 375 | Hyundai Getz | 2350000 | 2008 | |
| 407 | (Lada) 2114 () | 950000 | 2008 | |
| 421 | Chevrolet Captiva | 4800000 | 2008 | |
| 424 | Land Rover Range Rover | 11200000 | 2008 | - () |
| 439 | SsangYong Musso | 2400000 | 2008 | |
| 457 | Volkswagen Tiguan | 6000000 | 2008 | |
| 467 | (Lada) 2172 () | 900000 | 2008 | |
| 478 | Nissan Teana | 5000000 | 2008 | |
| 510 | (Lada) 2114 () | 890000 | 2008 | |
| 513 | (Lada) 2170 () | 1250000 | 2008 | |
| 527 | Toyota Land Cruiser Prado | 11800000 | 2008 | |
| 548 | Toyota Camry | 3999000 | 2008 | |
| 561 | Toyota Land Cruiser | 15700000 | 2008 | |
| 568 | Toyota Camry | 5500000 | 2008 | |
| 575 | Toyota Highlander | 10000000 | 2008 | |
| 576 | Audi Q7 | 5800000 | 2008 | - () |
| 582 | Toyota Land Cruiser | 12000000 | 2008 | |

| | Body | Engine_volume | Mileage | Transmission | Steering_wheel | \ |
|-----|------|---------------|---------|--------------|----------------|---|
| 7 | | 3.0 | 219000 | | | |
| 14 | | 3.5 | 210000 | | | |
| 41 | | 2.7 | 205000 | | | |
| 78 | | 2.9 | 395700 | | | |
| 151 | | 1.6 | 0 | | | |
| 182 | | 3.5 | 0 | | | |
| 185 | | 3.5 | 0 | | | |
| 206 | | 1.5 | 320 | | | |
| 217 | | 1.6 | 180000 | | | |
| 222 | | 3.0 | 139000 | | | |
| 225 | | 1.5 | 150000 | | | |
| 229 | | 1.5 | 190000 | | | |
| 238 | | 1.7 | 0 | | | |
| 250 | | 2.0 | 117000 | | | |

| | | |
|-----|-----|--------|
| 253 | 1.5 | 0 |
| 356 | 2.9 | 0 |
| 361 | 1.4 | 215000 |
| 375 | 1.4 | 160000 |
| 407 | 1.5 | 165000 |
| 421 | 2.4 | 141000 |
| 424 | 4.2 | 185000 |
| 439 | 3.2 | 0 |
| 457 | 2.0 | 167000 |
| 467 | 1.6 | 203630 |
| 478 | 3.5 | 199000 |
| 510 | 1.5 | 0 |
| 513 | 1.6 | 0 |
| 527 | 4.0 | 0 |
| 548 | 2.4 | 0 |
| 561 | 4.7 | 166000 |
| 568 | 2.4 | 0 |
| 575 | 3.5 | 0 |
| 576 | 4.2 | 110 |
| 582 | 4.7 | 212000 |

Color

7
 14
 41
 78
 151
 182
 185
 206
 217
 222
 225
 229
 238
 250
 253
 356
 361
 375
 407
 421
 424
 439
 457
 467
 478

510
513
527
548
561
568
575
576
582

```
[57]: import sklearn
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_auc_score, precision_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
```

0.0.3 Data preprocessing. Decision trees. We encode categorical features using an ordinal encoding scheme. Encode categorical features as a one-hot numeric array. `LabelEncoder` can be used to normalize labels. It can also be used to transform non-numerical labels (as long as they are hashable and comparable) to numerical labels.

```
[58]: data = df.drop('Price',axis='columns')
target = df['Price']
car_n = LabelEncoder()
body_n = LabelEncoder()
tr_n = LabelEncoder()
data['Car_2'] = car_n.fit_transform(data['Car'])
data['Body_2'] = body_n.fit_transform(data['Body'])
data['Transmission_2'] = tr_n.fit_transform(data['Transmission'])
data
```

```
[58]:
```

| | Car | Year | Region | Body \ |
|----|-------------------------|------|--------|--------|
| 0 | Chevrolet Tracker | 2015 | | |
| 1 | Opel Frontera | 1994 | | |
| 2 | Renault Sandero Stepway | 2014 | | |
| 3 | Toyota Land Cruiser | 2010 | | |
| 4 | Mitsubishi Pajero | 2007 | | |
| .. | ... | ... | ... | ... |

```

595          Daewoo Nexia  2007
596                      2015
597          BMW X5  2006  -  (  )
598  Volkswagen Transporter  1991
599  Mercedes-Benz E 230  1992  -  (  )

      Engine_volume  Mileage  Transmission  Steering_wheel  Color \
0          1.8      46000
1          2.3         0
2          1.6     110000
3          4.0     220000
4          3.0     137000
..          ...         ...         ...         ...         ...
595         1.5     190000
596         2.9         0
597         4.8     220000
598         2.0     222222
599         2.3        222

      Car_2  Body_2  Transmission_2
0         23      2              0
1        123      0              2
2        134     12              0
3        156      0              0
4        109      0              0
..          ...     ...         ...
595        30      8              2
596       192     11              2
597        12      2              0
598       174      6              2
599        86     10              2

```

[517 rows x 12 columns]

```
[59]: len(df.Car.unique())
```

```
[59]: 197
```

```
[60]: data = data[['Year', 'Engine_volume', 'Body_2']]
```

```
[61]: X, y = data, target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/3)
clf = LinearRegression()
clf.fit(X_train, y_train)
clf.score(X_test, y_test)
```

```
[61]: 0.6259556263293345
```

```
[62]: X_test
```

```
[62]:      Year  Engine_volume  Body_2
144  2005           2.4         8
62   2018           2.0         2
354  2019           5.6         0
248  2009           4.7         0
505  2005           1.6        12
..   ...           ...       ...
536  1989           2.3         8
179  2007           1.6        12
416  2017           2.0         2
485  1999           4.0         0
471  2010           1.6         8
```

```
[173 rows x 3 columns]
```

```
[63]: clf.predict(X_test)[:3]
```

```
[63]: array([ 4172022.42105854, 11208486.00488257, 19079302.39893699])
```

```
[64]: y_test[:3]
```

```
[64]: 144      4600000
62      11350000
354      29000000
Name: Price, dtype: int64
```

```
[65]: model = DecisionTreeClassifier()
model.fit(data, target)
model.score(data, target)
```

```
[65]: 0.7156673114119922
```

```
[66]: df.tail()
```

```
[66]:      Car      Price  Year      Region      Body \
595      Daewoo Nexia   950000  2007
596              6400000  2015
597      BMW X5   7000000  2006      -      (      )
598  Volkswagen Transporter  1600000  1991
599      Mercedes-Benz E 230  1800000  1992      -      (      )

      Engine_volume  Mileage  Transmission  Steering_wheel      Color
595              1.5    190000
596              2.9         0
597              4.8    220000
598              2.0    222222
```

```
[67]: model.predict([[2007,1.5,8]])
```

```
[67]: array([950000], dtype=int64)
```

```
[68]: param_grid = [
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001], 'kernel': ['rbf']},]
svc = svm.SVC()
grid_svc = GridSearchCV(svc, param_grid, cv=3)
grid_svc.fit(X_train, y_train)
```

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\model_selection_split.py:657: Warning: The least populated class in y has only 1 members, which is too few. The minimum number of members in any class cannot be less than n_splits=3.

% (min_groups, self.n_splits)), Warning)

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\model_selection_search.py:814: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.

DeprecationWarning)

```
[68]: GridSearchCV(cv=3, error_score='raise-deprecating',
    estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3,
    gamma='auto_deprecated', kernel='rbf', max_iter=-1,
    probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False),
    iid='warn', n_jobs=None,
    param_grid=[{'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001],
    'kernel': ['rbf']}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)
```

```
[70]: clf_svc = grid_svc.best_estimator_
print(clf_svc)
```

```
SVC(C=10, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
[71]: param_grid = [{'n_neighbors': [2,3,4,5,6]}]
knn = KNeighborsClassifier()
```

```
grid_knn = GridSearchCV(knn, param_grid, cv=3)
grid_knn.fit(X_train, y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_split.py:657: Warning: The least populated class in y has only 1 members, which is too few. The minimum number of members in any class cannot be less than n_splits=3.

```
% (min_groups, self.n_splits)), Warning)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_search.py:814: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.

```
DeprecationWarning)
```

```
[71]: GridSearchCV(cv=3, error_score='raise-deprecating',
                  estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                metric='minkowski',
                                                metric_params=None, n_jobs=None,
                                                n_neighbors=5, p=2,
                                                weights='uniform'),
                  iid='warn', n_jobs=None,
                  param_grid=[{'n_neighbors': [2, 3, 4, 5, 6]}],
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)
```

```
[72]: clf_knn = grid_knn.best_estimator_
       print(clf_knn)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=2, p=2,
                    weights='uniform')
```

```
[73]: decision_tree = DecisionTreeClassifier()
       decision_tree.fit(X_train, y_train)
       for model in [clf_knn, clf_svc, decision_tree]:
           y_pred = model.predict(X_test)
           score = precision_score(y_test.to_numpy().T, y_pred, average='weighted')
           print(model, 'score:', score)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=2, p=2,
                    weights='uniform') score: 0.03177814478392513
```

```
SVC(C=10, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False) score: 0.001707238866116075
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                      max_features=None, max_leaf_nodes=None,
```

```
min_impurity_decrease=0.0, min_impurity_split=None,  
min_samples_leaf=1, min_samples_split=2,  
min_weight_fraction_leaf=0.0, presort=False,  
random_state=None, splitter='best') score:
```

0.024116891457931923

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples.

```
'precision', 'predicted', average, warn_for)
```

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples.


```
'precision', 'predicted', average, warn_for)
```

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples.

```
'precision', 'predicted', average, warn_for)
```

```
[74]: df = pd.read_csv('kolesa.csv')  
df = df[:40]  
x = df['Year']  
y = df['Car']  
plt.bar(x, y)  
  
plt.title("Simple plot 1")  
plt.ylabel("y-axis")  
plt.xlabel("x-axis")  
  
plt.show()
```



```
[75]: df = df.groupby(['Engine_volume']).first()
df
```

```
[75]:
```

| | Unnamed: 0 | Car | Price | Year | \ |
|---------------|------------|-------------------------|----------|------|---|
| Engine_volume | | | | | |
| 1.6 | 2 | Renault Sandero Stepway | 3990000 | 2014 | |
| 1.8 | 0 | Chevrolet Tracker | 5850000 | 2015 | |
| 2.0 | 8 | Mitsubishi RVR | 1600000 | 1995 | |
| 2.2 | 10 | Opel Zafira | 2300000 | 2001 | |
| 2.3 | 1 | Opel Frontera | 1650000 | 1994 | |
| 2.4 | 5 | Toyota Camry | 3800000 | 2003 | |
| 2.5 | 6 | Mazda MPV | 3000000 | 1999 | |
| 2.8 | 9 | Audi 100 | 2500000 | 1994 | |
| 3.0 | 4 | Mitsubishi Pajero | 8000000 | 2007 | |
| 3.5 | 14 | Toyota Highlander | 8800000 | 2008 | |
| 4.0 | 3 | Toyota Land Cruiser | 16800000 | 2010 | |
| 4.4 | 19 | Land Rover Range Rover | 3000000 | 2004 | |
| 8.0 | 15 | Daewoo Matiz | 900000 | 2000 | |

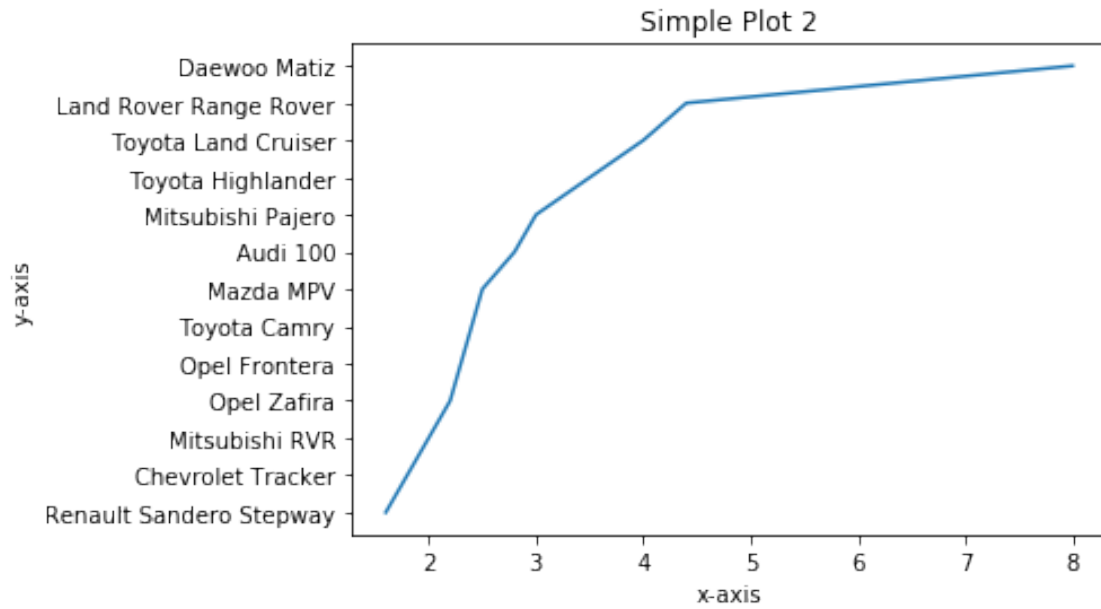
| | Region | Body | Mileage | Transmission | \ |
|---------------|--------|------|---------|--------------|---|
| Engine_volume | | | | | |
| 1.6 | | | 110000 | | |
| 1.8 | | | 46000 | | |
| 2.0 | | | 444444 | | |
| 2.2 | | | 160000 | | |
| 2.3 | | | 0 | | |
| 2.4 | | | 0 | | |

| | |
|-----|--------|
| 2.5 | 0 |
| 2.8 | 300000 |
| 3.0 | 137000 |
| 3.5 | 210000 |
| 4.0 | 220000 |
| 4.4 | 0 |
| 8.0 | 0 |

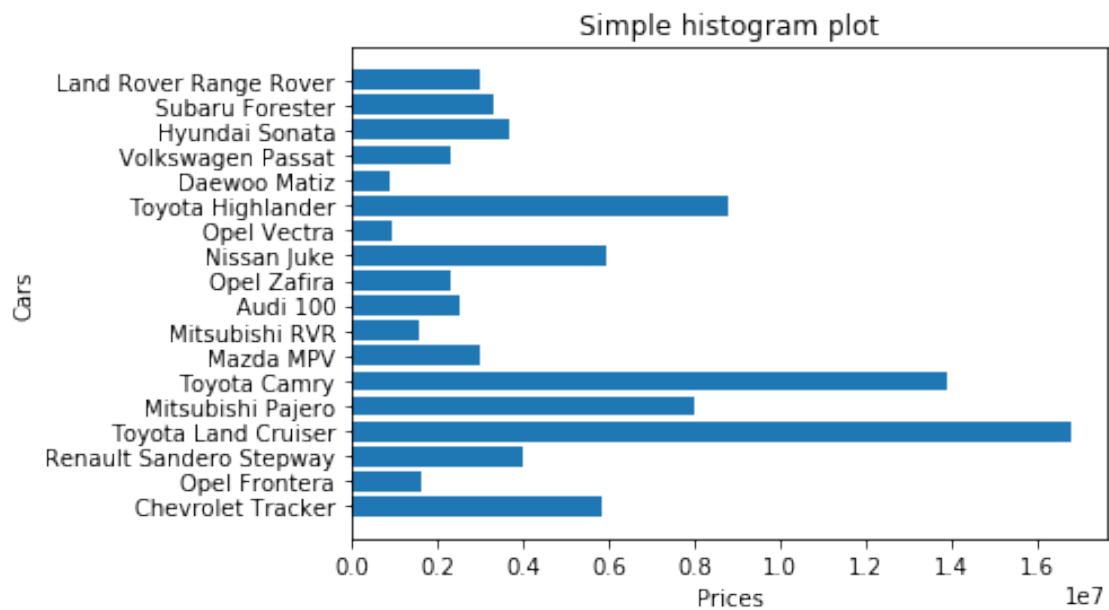
| Engine_volume | Steering_wheel | Color |
|---------------|----------------|-------|
| 1.6 | | |
| 1.8 | | |
| 2.0 | | |
| 2.2 | | |
| 2.3 | | |
| 2.4 | | |
| 2.5 | | |
| 2.8 | | |
| 3.0 | | |
| 3.5 | | |
| 4.0 | | |
| 4.4 | | |
| 8.0 | | |

```
[76]: x = df['Car']
      y = df.index
      plt.plot(y,x)
      plt.title("Simple Plot 2")
      plt.ylabel("y-axis")
      plt.xlabel("x-axis")

      plt.show()
```

```
[77]: df = pd.read_csv('kolesa.csv')
df = df[:20]
plt.barh(df['Car'], df['Price'])
plt.title("Simple histogram plot")
plt.xlabel('Prices')
plt.ylabel('Cars')
plt.show()
```



```
[78]: df = pd.read_csv('kolesa.csv')
df = df.groupby(['Body']).first()
df[2:13]
```

```
[78]:
```

| | Unnamed: 0 | Car | Price | Year | \ |
|------|------------|------------------------|----------|------|---|
| Body | | | | | |
| | 1 | Opel Frontera | 1650000 | 1994 | |
| | 244 | BMW 650 | 10000000 | 2007 | |
| | 0 | Chevrolet Tracker | 5850000 | 2015 | |
| | 250 | Hyundai Tuscani | 3100000 | 2008 | |
| | 176 | Toyota Carina E | 1900000 | 1992 | |
| | 86 | Volkswagen Transporter | 1800000 | 1997 | |
| | 6 | Mazda MPV | 3000000 | 1999 | |
| | 51 | Toyota Tundra | 23000000 | 2017 | |
| | 5 | Toyota Camry | 3800000 | 2003 | |
| | 587 | Toyota Supra | 3000000 | 1991 | |
| | 9 | Audi 100 | 2500000 | 1994 | |

| | Region | Engine_volume | Mileage | Transmission | \ |
|------|--------|---------------|---------|--------------|---|
| Body | | | | | |
| | | 2.3 | 0 | | |
| | | 4.8 | 137000 | | |
| | | 1.8 | 46000 | | |
| | | 2.0 | 117000 | | |
| | | 2.0 | 0 | | |
| | | 1.6 | 0 | | |
| | | 2.5 | 0 | | |
| | | 4.6 | 0 | | |
| | | 2.4 | 0 | | |
| | - () | 2.5 | 0 | | |
| | | 2.8 | 300000 | | |

| | Steering_wheel | Color |
|------|----------------|-------|
| Body | | |

```
[25]: import numpy as np
x = df.index[2:13]
y = df['Region'][2:13]

plt.ylabel("Body")
plt.xlabel("Region")
plt.title("Simple Scatter plot")
rng = np.random.RandomState(0)
x = df.index[2:10]
y = df['Region'][2:10]
colors = rng.rand(8)
sizes = 1000 * rng.rand(8)
plt.scatter(y, x, c=colors, s=sizes)
plt.colorbar()
```

[25]: <matplotlib.colorbar.Colorbar at 0x259e984b548>

